

Final Project Code

```
In [ ]: #load libraries
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import missingno
from mpl_toolkits.mplot3d import Axes3D
from sklearn.feature_extraction.text import CountVectorizer
import nltk
from nltk.corpus import stopwords
from nltk.stem.porter import PorterStemmer
import collections
from collections import Counter
import string
import re
from wordcloud import WordCloud
from nltk.tokenize import word_tokenize
from nltk.probability import FreqDist
from nltk import bigrams
from pprint import pprint
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
from sklearn.ensemble import RandomForestClassifier
from sklearn import tree
from sklearn.model_selection import train_test_split, RandomizedSearchCV, GridSearchCV
from sklearn.preprocessing import PowerTransformer, OneHotEncoder
from sklearn.compose import ColumnTransformer
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from keras.models import Sequential
from keras.layers import Dense, Embedding, LSTM, SpatialDropout1D, GRU, Dropout
from tensorflow.keras.utils import to_categorical
from keras.callbacks import EarlyStopping, ModelCheckpoint

# View all columns and rows
pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', 90) # set to fit 90 column descriptions

# Set up notebook to display multiple outputs in one cell
from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"
```

```
In [ ]: tweets_df = pd.read_csv('/content/tweets.csv')
```

EDA

```
In [ ]: # Inspect dataset

## Inspect columns and variable types
tweets_df.info()
```

```
## View dataframe rows
tweets_df.head()

## View summary statistics
tweets_df.describe()

## View missing values
tweets_df.isnull().sum()

## Plot missing values
missingno.matrix(tweets_df)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14640 entries, 0 to 14639
Data columns (total 15 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   tweet_id                             14640 non-null  int64
 1   airline_sentiment                    14640 non-null  object
 2   airline_sentiment_confidence         14640 non-null  float64
 3   negativereason                       9178 non-null   object
 4   negativereason_confidence            10522 non-null  float64
 5   airline                              14640 non-null  object
 6   airline_sentiment_gold                40 non-null     object
 7   name                                 14640 non-null  object
 8   negativereason_gold                  32 non-null     object
 9   retweet_count                        14640 non-null  int64
10  text                                 14640 non-null  object
11  tweet_coord                           1019 non-null   object
12  tweet_created                         14640 non-null  object
13  tweet_location                        9907 non-null   object
14  user_timezone                         9820 non-null   object
dtypes: float64(2), int64(2), object(11)
memory usage: 1.7+ MB
```

Out[]: tweet_id airline_sentiment airline_sentiment_confidence negativereason negati

0	570306133677760513	neutral	1.0000	NaN
---	--------------------	---------	--------	-----

1	570301130888122368	positive	0.3486	NaN
---	--------------------	----------	--------	-----

2	570301083672813571	neutral	0.6837	NaN
---	--------------------	---------	--------	-----

3	570301031407624196	negative	1.0000	Bad Flight
---	--------------------	----------	--------	------------


4	570300817074462722	negative	1.0000	Can't Tell
---	--------------------	----------	--------	------------

	tweet_id	airline_sentiment_confidence	negativereason_confidence	retweet_count
count	1.464000e+04	14640.000000	10522.000000	14640.000000
mean	5.692184e+17	0.900169	0.638298	0.082650
std	7.791112e+14	0.162830	0.330440	0.745778
min	5.675883e+17	0.335000	0.000000	0.000000
25%	5.685592e+17	0.692300	0.360600	0.000000
50%	5.694779e+17	1.000000	0.670600	0.000000
75%	5.698905e+17	1.000000	1.000000	0.000000
max	5.703106e+17	1.000000	1.000000	44.000000

```

Out[ ]: tweet_id                0
airline_sentiment            0
airline_sentiment_confidence 0
negativereason              5462
negativereason_confidence   4118
airline                     0
airline_sentiment_gold      14600
name                        0
negativereason_gold         14608
retweet_count               0
text                       0
tweet_coord                 13621
tweet_created               0
tweet_location              4733
user_timezone               4820
dtype: int64
Out[ ]: <Axes: >

```



```

In [ ]: # Convert airline_sentiment classification to numeric: positive, neutral, negative
tweets_df['numeric_sentiment'] = tweets_df.apply(lambda x: 1 if (x['airline_sentiment'] == 'positive') else 0 if (x['airline_sentiment'] == 'neutral') else -1 if (x['airline_sentiment'] == 'negative') else 0)

# Create weighted_sentiment feature, the product of the sentiment and the confidence
tweets_df['weighted_sentiment'] = tweets_df['numeric_sentiment'] * tweets_df['airline_sentiment_confidence']

```

```
In [ ]: # Compute summary statistics for numeric_sentiment variable.

print("Summary stats for numeric_sentiment.")
tweets_df['numeric_sentiment'].describe()

print("Median and mode for numeric_sentiment.")
tweets_df['numeric_sentiment'].median()
tweets_df['numeric_sentiment'].mode().iloc[0] # Taking first mode if multiple.

print("Summary stats for weighted_sentiment.")
tweets_df['weighted_sentiment'].describe()

print("Median and mode for weighted_sentiment.")
tweets_df['weighted_sentiment'].median()
tweets_df['weighted_sentiment'].mode().iloc[0] # Taking first mode if multiple.

## Noting that the 'numeric_sentiment' and 'weighted_sentiment' columns are pre
```

```
Summary stats for numeric_sentiment.
Out[ ]: count    14640.000000
        mean      -0.465505
        std       0.756084
        min      -1.000000
        25%      -1.000000
        50%      -1.000000
        75%       0.000000
        max       1.000000
        Name: numeric_sentiment, dtype: float64
        Median and mode for numeric_sentiment.
```

```
Out[ ]: -1.0
```

```
Out[ ]: -1
```

```
Summary stats for weighted_sentiment.
Out[ ]: count    14640.000000
        mean      -0.444385
        std       0.698999
        min      -1.000000
        25%      -1.000000
        50%      -1.000000
        75%       0.000000
        max       1.000000
        Name: weighted_sentiment, dtype: float64
        Median and mode for weighted_sentiment.
```










```
Out[ ]: -1.0
```

```
Out[ ]: -1.0
```

```
In [ ]: # look at Airline Sentiment distribution
tweets_df.hist(column='airline', by='numeric_sentiment',figsize=(12,8))
tweets_df.hist(column='weighted_sentiment', by='airline',figsize=(12,8))
#sns.histplot(data_df,x='airline',hue='num_sentiment')
#plt.show()
```

```
Out[ ]: array([[<Axes: title={'center': '-1'}>, <Axes: title={'center': '0'}>],
               [<Axes: title={'center': '1'}>, <Axes: >]], dtype=object)
```


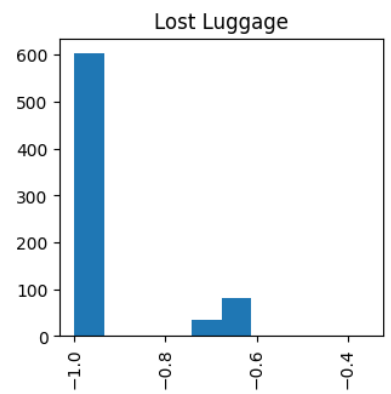
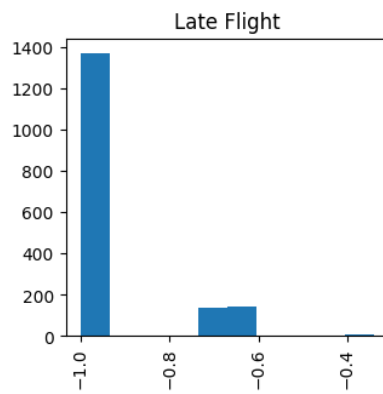
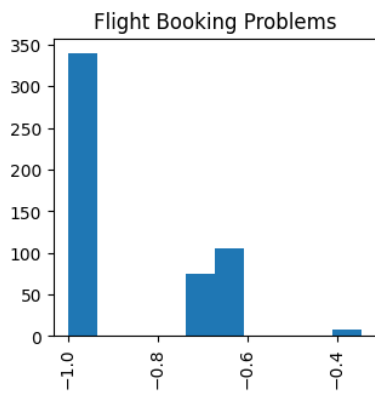
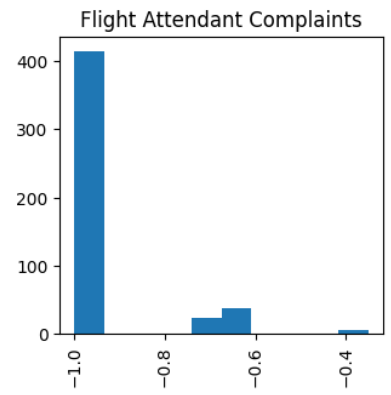
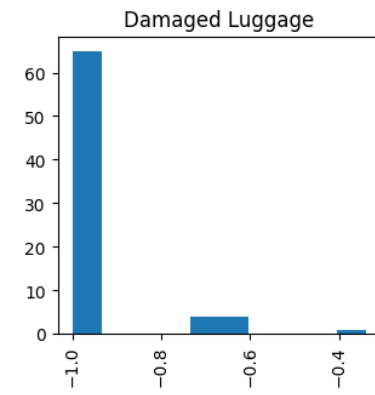
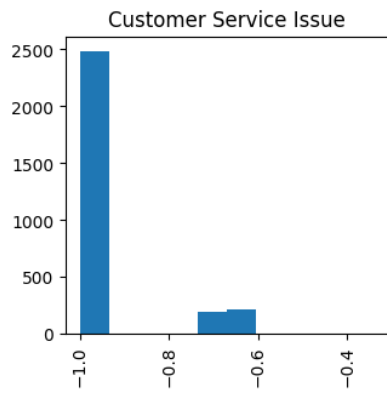
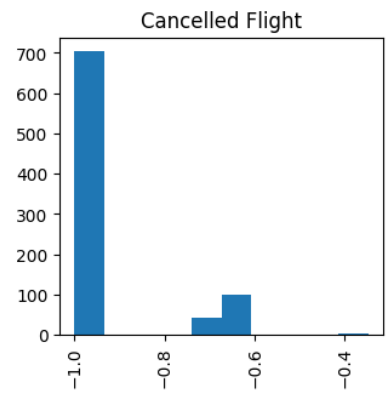
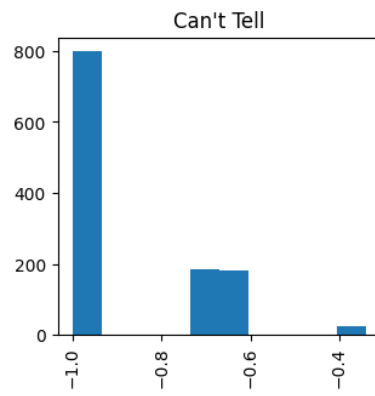
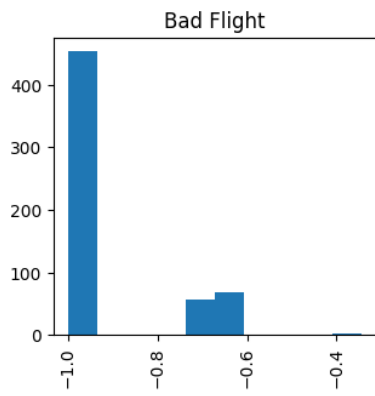
```
Out[ ]: array([[<Axes: title={'center': 'American'}>,
<Axes: title={'center': 'Delta'}>],
[<Axes: title={'center': 'Southwest'}>,
<Axes: title={'center': 'US Airways'}>],
[<Axes: title={'center': 'United'}>,
<Axes: title={'center': 'Virgin America'}>]], dtype=object)
```

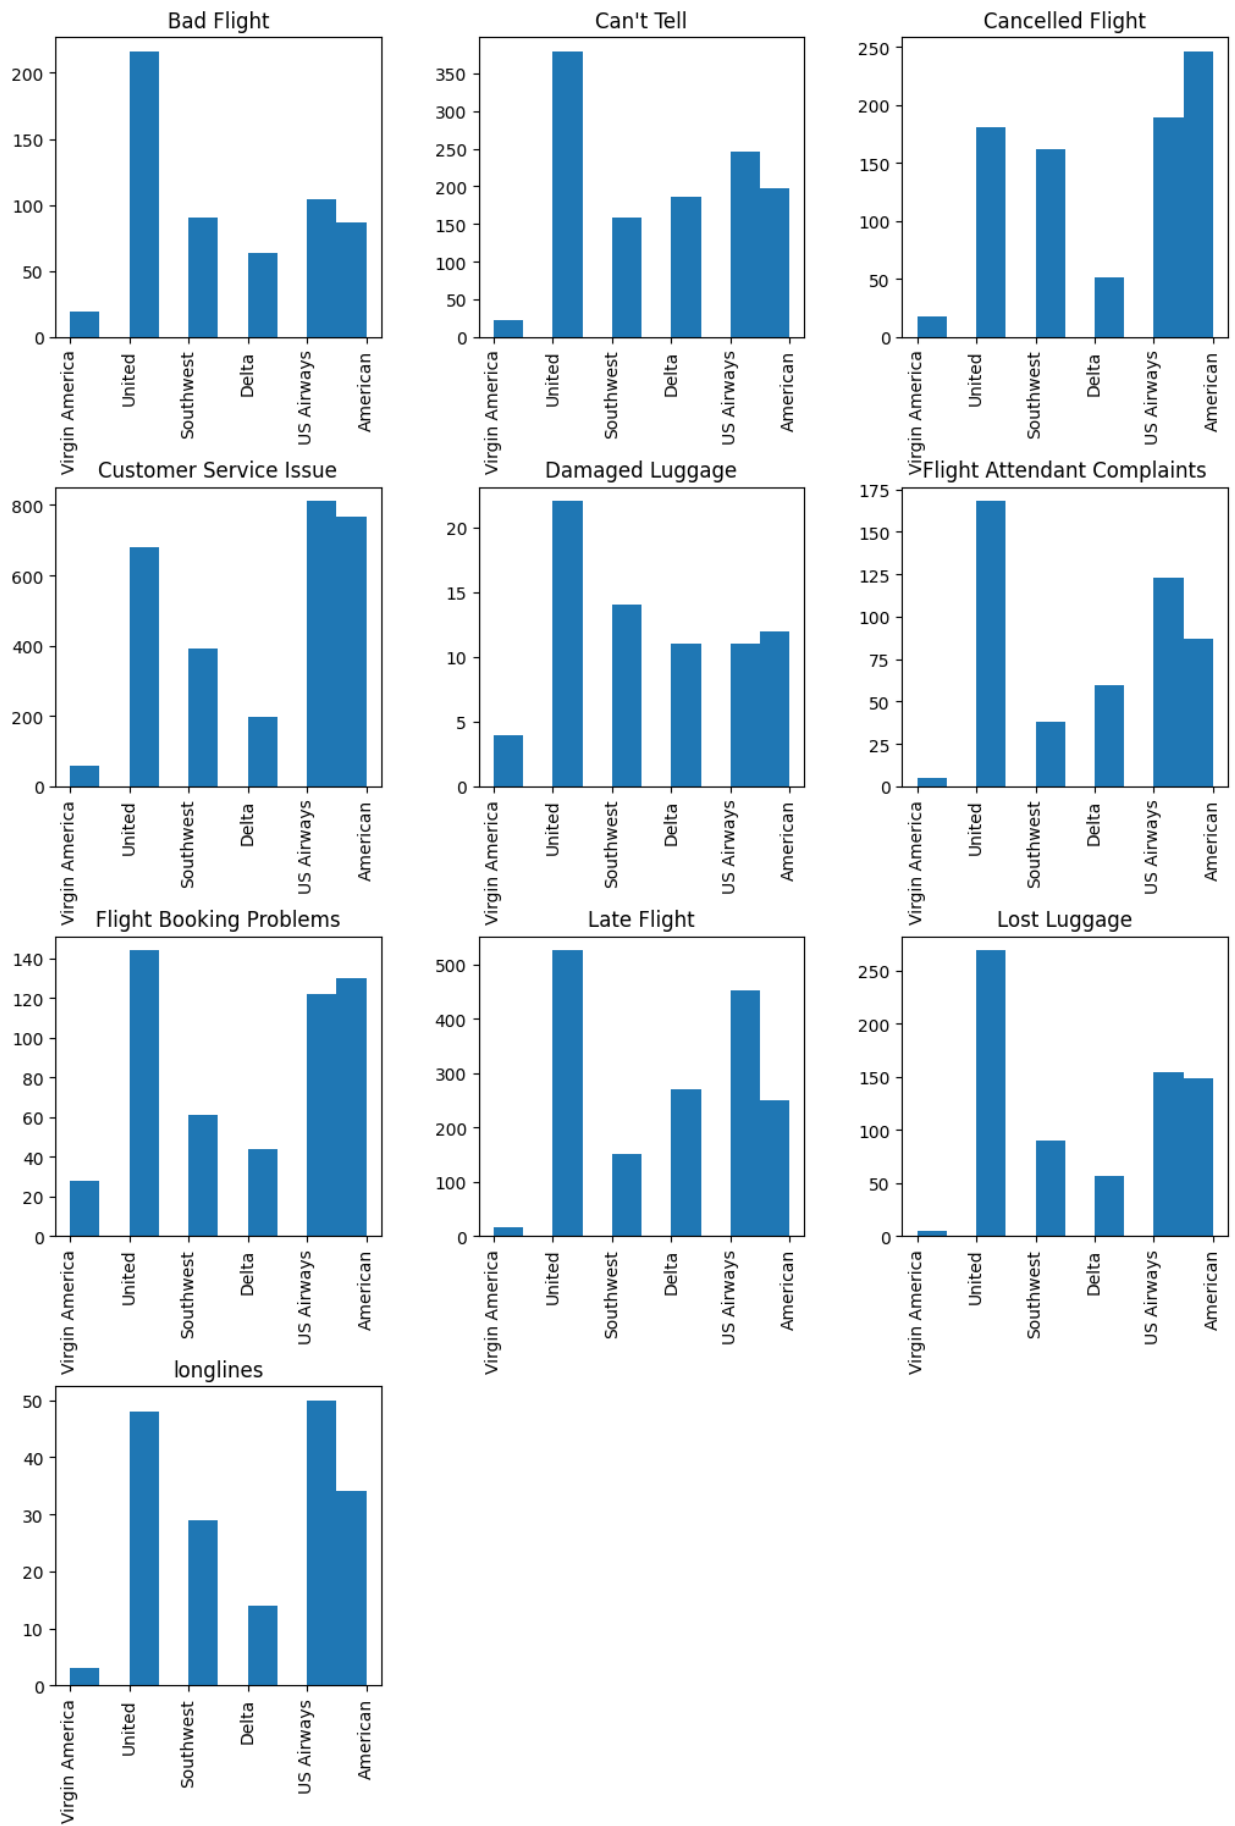


```
In [ ]: # look at Negative Reason distribution
data_neg = tweets_df[tweets_df['numeric_sentiment'] < 0]
data_neg.hist(column='weighted_sentiment', by='negativereason', figsize=(12,18))
data_neg.hist(column='airline', by='negativereason', figsize=(12,18))
```

```
Out[ ]: array([[<Axes: title={'center': 'Bad Flight'}>,
<Axes: title={'center': 'Can't Tell'}>,
<Axes: title={'center': 'Cancelled Flight'}>],
[<Axes: title={'center': 'Customer Service Issue'}>,
<Axes: title={'center': 'Damaged Luggage'}>,
<Axes: title={'center': 'Flight Attendant Complaints'}>],
[<Axes: title={'center': 'Flight Booking Problems'}>,
<Axes: title={'center': 'Late Flight'}>,
<Axes: title={'center': 'Lost Luggage'}>],
[<Axes: title={'center': 'longlines'}>, <Axes: >, <Axes: >]],
dtype=object)
```

```
Out[ ]: array([[<Axes: title={'center': 'Bad Flight'}>,
<Axes: title={'center': 'Can't Tell'}>,
<Axes: title={'center': 'Cancelled Flight'}>],
[<Axes: title={'center': 'Customer Service Issue'}>,
<Axes: title={'center': 'Damaged Luggage'}>,
<Axes: title={'center': 'Flight Attendant Complaints'}>],
[<Axes: title={'center': 'Flight Booking Problems'}>,
<Axes: title={'center': 'Late Flight'}>,
<Axes: title={'center': 'Lost Luggage'}>],
[<Axes: title={'center': 'longlines'}>, <Axes: >, <Axes: >]],
dtype=object)
```





In []: `#create visuals`


```

nltk.download('stopwords')
nltk.download('punkt')

# Check for missing values
print(tweets_df.isnull().sum())

# Explore sentiment distribution
sns.countplot(x='airline_sentiment', data=tweets_df, order=tweets_df['airline_s
plt.title('Sentiment Distribution')
plt.show()

# Explore sentiment distribution for each airline
sns.countplot(x='airline', hue='airline_sentiment', data=tweets_df)
plt.title('Sentiment Distribution by Airline')
plt.show()

# Word cloud for positive, negative, and neutral sentiments
company_names = ['Virgin America', 'United', 'Southwest', 'Delta', 'US Airways',
                  'Southwest Air', 'SouthwestAir', 'VirginAmerica', 'AmericanAir
stop_words = set(stopwords.words('english') + company_names)

def plot_wordcloud(sentiment):
    words = ' '.join(tweets_df[tweets_df['airline_sentiment'] == sentiment]['text'])
    wordcloud = WordCloud(width=800, height=400, stopwords = stop_words, random_state=1)
    plt.figure(figsize=(10, 7))
    plt.imshow(wordcloud, interpolation="bilinear")
    plt.axis('off')
    plt.title(f'Word Cloud for {sentiment} Sentiment')
    plt.show()

plot_wordcloud('positive')
plot_wordcloud('negative')
plot_wordcloud('neutral')

# Tokenize and analyze word frequencies

def process_text(text):
    words = word_tokenize(text)
    words = [word.lower() for word in words if word.isalpha() and word.lower() != 'a']
    return words

tweets_df['processed_text'] = tweets_df['text'].apply(process_text)

# Calculate and plot word frequencies
all_words = [word for sublist in tweets_df['processed_text'] for word in sublist]
freq_dist = FreqDist(all_words)
freq_dist.plot(30, cumulative=False)

```

```

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!

```

Out[]: True

```

[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!


```

Out[]: True

tweet_id	0
airline_sentiment	0
airline_sentiment_confidence	0
negativereason	5462
negativereason_confidence	4118
airline	0
airline_sentiment_gold	14600
name	0
negativereason_gold	14608
retweet_count	0
text	0
tweet_coord	13621
tweet_created	0
tweet_location	4733
user_timezone	4820
numeric_sentiment	0
weighted_sentiment	0
dtype:	int64

Out[]: <Axes: xlabel='airline_sentiment', ylabel='count'>

Out[]: Text(0.5, 1.0, 'Sentiment Distribution')



Out[]: <Axes: xlabel='airline', ylabel='count'>


Out[]: Text(0.5, 1.0, 'Sentiment Distribution by Airline')

A grouped bar chart showing the count of sentiment (neutral, positive, negative) for six airlines: Virgin America, United, Southwest, Delta, US Airways, and American. The y-axis is labeled 'count' and ranges from 0 to 2500. The legend indicates: neutral (blue), positive (orange), and negative (green).

airline	neutral	positive	negative
Virgin America	170	160	180
United	700	500	2650
Southwest	670	570	1190
Delta	730	550	960
US Airways	390	270	2270
American	470	340	1960

[illegible]

[illegible][illegible]



Out[]: <Axes: xlabel='Samples', ylabel='Counts'>

Cleaning the Data

```
In [ ]: # We note that the majority of airline_sentiment_gold is empty. However, non-null
diffs = tweets_df[tweets_df['airline_sentiment_gold'].notnull()]
diffs = (diffs['airline_sentiment'] != diffs['airline_sentiment_gold']).any()
print(diffs)

## There are no differences between airline_sentiment and airline_sentiment_gold
tweets_df.drop(columns=['airline_sentiment_gold'], inplace=True)

False
```

```
In [ ]: # However, there are differences between those of negativereason and negativereason_gold

# If confidence level in airline_sentiment_confidence is higher on average we'd expect
sentiment_mean = tweets_df.groupby('negativereason')['airline_sentiment_confidence'].mean()
gold_sent_mean = tweets_df.groupby('negativereason_gold')['airline_sentiment_confidence'].mean()

print(f"Mean of confidence for negativereason:\n{sentiment_mean}")
print(f"\nMean of confidence for negativereason_gold:\n{gold_sent_mean}")

## Since negativereason_gold has a higher mean confidence level in the value of
```

```
tweets_df['imputed_neg_reason'] = tweets_df['negativereason_gold'].fillna(tweets_df['negativereason'])
check_df = tweets_df[(tweets_df['negativereason_gold'].notnull()) & (tweets_df['negativereason'].notnull())]
check_df = check_df[['negativereason', 'negativereason_gold', 'text']]
check_df

## A spot check confirms negativereason_gold offers more adequate information
```

Mean of confidence for negativereason:

negativereason	
Bad Flight	0.925816
Can't Tell	0.885067
Cancelled Flight	0.941773
Customer Service Issue	0.951616
Damaged Luggage	0.955727
Flight Attendant Complaints	0.951697
Flight Booking Problems	0.876460
Late Flight	0.940308
Lost Luggage	0.944268
longlines	0.943487

Name: airline_sentiment_confidence, dtype: float64

Mean of confidence for negativereason_gold:

negativereason_gold	
Bad Flight	0.965800
Can't Tell	0.913100
Cancelled Flight	1.000000
Cancelled Flight\nCustomer Service Issue	0.924850
Customer Service Issue	0.990667
Customer Service Issue\nCan't Tell	0.799100
Customer Service Issue\nLost Luggage	1.000000
Flight Attendant Complaints	0.955300
Late Flight	0.927025
Late Flight\nCancelled Flight	1.000000
Late Flight\nFlight Attendant Complaints	1.000000
Late Flight\nLost Luggage	1.000000
Lost Luggage\nDamaged Luggage	1.000000

Name: airline_sentiment_confidence, dtype: float64

Out []:	negativereason	negativereason_gold	text
1286	Late Flight	Late Flight\nFlight Attendant Complaints	@united I'm aware of the flight details, thank...
2017	Late Flight	Late Flight\nLost Luggage	@united flighted delayed for hours. 10pm arriv...
3149	Customer Service Issue	Cancelled Flight\nCustomer Service Issue	@united rebooked 24 hours after original fligh...
6530	Customer Service Issue	Cancelled Flight\nCustomer Service Issue	@SouthwestAir I never got a Cancelled Flightla...
8536	Lost Luggage	Lost Luggage\nDamaged Luggage	@JetBlue I am heading to JFK now just on princ...
12025	Cancelled Flight	Late Flight\nCancelled Flight	@AmericanAir over the last year 50% of my flig...
12034	Lost Luggage	Customer Service Issue\nLost Luggage	@AmericanAir you need to work harder on the di...
12038	longlines	Customer Service Issue	@DeltaAssist now at 57 minutes waiting on Silv...
12039	Can't Tell	Customer Service Issue\nCan't Tell	@DeltaAssist what I have to say is more than 1...

Feature Engineering

```
In [ ]: # Cleaning user_timezone values in stages. tweet_location not as useful (e.g.,
# Clean timezone column
tweets_df['clean_timezone'] = tweets_df['user_timezone'].where(tweets_df['user_

# Fill NaN values in the original column
tweets_df['user_timezone'].fillna('N/A', inplace=True)

# Unclean timezone column
tweets_df['unclean_timezone'] = tweets_df['user_timezone'].where(~tweets_df['c

## Code leaves in unclean_timezone column. We could delete the above line and s
```

```
In [ ]: # Convert 'tweet_created' column to datetime dtype
tweets_df['tweet_created'] = pd.to_datetime(tweets_df['tweet_created'])

# Extract day of the month
tweets_df['tweet_day'] = tweets_df['tweet_created'].dt.day

# Extract time of day
tweets_df['time_of_day'] = pd.cut(
    tweets_df['tweet_created'].dt.hour,
    bins=[0, 6, 12, 18, 24],
    labels=['Night', 'Morning', 'Afternoon', 'Evening'],
    include_lowest=True
)
```

Pull Only Data Columns of Interest

```
In [ ]: #clean up data set
data_pre_process=tweets_df[['text','numeric_sentiment', 'weighted_sentiment'],'a
```

Initial Cleaning of Text Column and Prepare Data for Thorough Cleaning

```
In [ ]: #clean the text column for random forest modelling to identify key word counts
from bs4 import BeautifulSoup
import re
import nltk
# nltk.download()
from nltk.corpus import stopwords # Import the stop word list
nltk.download('stopwords')

data_size = (data_pre_process['text'].size)
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
Out[ ]: True
```

Cleaning the Data Using Beautiful Soup

```
In [ ]: #function to clean text usint Beautiful Soup
def clean_text_data(data_point, data_size):
    review_soup = BeautifulSoup(data_point)
    review_text = review_soup.get_text()
    review_letters_only = re.sub("[^a-zA-Z]", " ", review_text)
    review_lower_case = review_letters_only.lower()
    review_words = review_lower_case.split()
    stop_words = stopwords.words("english")
    meaningful_words = [x for x in review_words if x not in stop_words]

    if (i%2000 == 0 ):
        print("Cleaned %d of %d data (%d %%)."% ( i, data_size, ((i)/data_size))

    return( " ".join( meaningful_words))

for i in range(data_size):
    data_pre_process["text"][i] = clean_text_data(data_pre_process["text"][i],
    print("Cleaning training completed!")
```

```
<ipython-input-15-85e9262fa24a>:17: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
data_pre_process["text"][i] = clean_text_data(data_pre_process["text"][i], d
ata_size)
```

```
<ipython-input-15-85e9262fa24a>:3: MarkupResemblesLocatorWarning: The input lo
oks more like a filename than markup. You may want to open this file and pass
the filehandle into BeautifulSoup.
```

```
review_soup = BeautifulSoup(data_point)
```



```
Cleaned 0 of 14640 data (0 %).
Cleaned 2000 of 14640 data (13 %).
Cleaned 4000 of 14640 data (27 %).
Cleaned 6000 of 14640 data (40 %).
Cleaned 8000 of 14640 data (54 %).
Cleaned 10000 of 14640 data (68 %).
Cleaned 12000 of 14640 data (81 %).
Cleaned 14000 of 14640 data (95 %).
Cleaning training completed!
```

Executing a Random Forest Model Without Fitting to Learn Top Phrases and Number of Occurences for Each

```
In [ ]: vectorizer = CountVectorizer(analyzer = "word", \
                                   ngram_range= (2,3), \
                                   tokenizer = None, \
                                   preprocessor = None, \
                                   stop_words = ['united','usairways','americanair',

X_train, X_cv, Y_train, Y_cv = train_test_split(data_pre_process["text"], data

X_train = vectorizer.fit_transform(X_train)
X_train = X_train.toarray()
print(X_train.shape)

X_cv = vectorizer.transform(X_cv)
X_cv = X_cv.toarray()
print(X_cv.shape)

vocab = vectorizer.get_feature_names_out()
distribution = np.sum(X_train, axis=0)


#for tag, count in zip(distribution[:100],vocab[:100]):
#    print(count, tag)

zipped = zip (distribution,vocab)
zipped = list(zipped)
zipped.sort(reverse=True)
zip_top = zipped[:10]
zip_top_df = pd.DataFrame(zip_top)
#zip_top_df.head()
graph = sns.barplot(zip_top_df,x=zip_top_df[1],y=zip_top_df[0])
graph.set_xticklabels(
    labels=zip_top_df[1], rotation=-30)
plt.show()

(11712, 141166)
(2928, 141166)

<ipython-input-16-8eb1748bbf50>:31: UserWarning: FixedFormatter should only be
used together with FixedLocator
graph.set_xticklabels(
```

```
Out[ ]: [Text(0, 0, 'customer service'),
Text(1, 0, 'cancelled flightled'),
Text(2, 0, 'late flight'),
Text(3, 0, 'flight cancelled'),
Text(4, 0, 'cancelled flighted'),
Text(5, 0, 'fleet fleek'),
Text(6, 0, 'booking problems'),
Text(7, 0, 'flight cancelled flightled'),
Text(8, 0, 'late flightr'),
Text(9, 0, 'cancelled flight')]
```



Bag of Words

```
In [ ]: tweets_df1 = tweets_df.copy()
tweets_df1['text'] = tweets_df1['text'].astype(str)
```

```
In [ ]: stop_words = set(stopwords.words('english'))
stop_words.remove('not') #need 'not' as a negative identifier in tweets
string.punctuation
```

```
Out[ ]: '!"#$%&\'()*+,-./:;<=>?@[\\]^_`{|}~'
```

```
In [ ]: #clean Tweet text
ps = PorterStemmer()
```

```

tweet_list = []

for i in range(0, 14640):
    tweet = re.sub('@[\s]+', '', tweets_df1['text'][i])
    tweet = re.sub('http[\s]+', '', tweet)
    tweet = re.sub('[\sstring.punctuation+]', '', tweet)
    tweet = tweet.lower()
    tweet = tweet.split()
    tweet = [ps.stem(word) for word in tweet if not word in stop_words]
    tweet = ' '.join(tweet)
    tweet_list.append(tweet)

```

```

In [ ]: #remove emojis
def remove_emoji(text):
    emojis = re.compile("[
        u\"\\U0001F600-\\U0001F64F\" # emoticons
        u\"\\U0001F300-\\U0001F5FF\" # symbols & pictographs
        u\"\\U0001F680-\\U0001F6FF\" # transport & map symbols
        u\"\\U0001F1E0-\\U0001F1FF\" # flags (iOS)
        u\"\\U00002702-\\U000027B0\"
        u\"\\U000024C2-\\U0001F251\"
    ]+", flags=re.UNICODE)

    text = re.sub(emojis, '', text)
    return text

tweet_list = list(map(remove_emoji, tweet_list))

```

```

In [ ]: len(tweet_list)
tweet_list[1:5]

```

```

Out[ ]: 14640

```

```

Out[ ]: ['plu youv ad commerci experi tacki',
        'didnt today must mean need take anoth trip',
        'realli aggress blast obnoxii entertain guest face amp littl recours',
        'realli big bad thing']

```

```

In [ ]: #find unique words for count vectorizer
text = ' '.join(tweet_list)
words = text.split()
unique_words = set(words)
print(len(unique_words))

#look at word counts
word_count = Counter(words)
sorted_word_count = sorted(word_count.items(), key = lambda x: x[1], reverse=True)
for word, count in sorted_word_count:
    print(f'{word}: {count}')

#find all words that occur at least 5 times
count_mult_words = sum(1 for word, count in sorted_word_count if count >= 5)
print(count_mult_words)

```

Streaming output truncated to the last 5000 lines.

ridic: 1
paycheck: 1
thingunit: 1
crackersnabisco: 1
pepperidg: 1
farm: 1
2daysofhel: 1
exitrow: 1
rj145: 1
lta: 1
answerthi: 1
whack: 1
circul: 1
burningman: 1
mkwlkr: 1
b737900: 1
9may5sep: 1
ua6465: 1
lostyou: 1
63446373: 1
sfodfw: 1
1576: 1
sfoord: 1
disappointedunit: 1
waspaid: 1
ua6255: 1
winnipeg: 1
ua1059: 1
tightconnect: 1
coltsmissingbag: 1
lightyear: 1
frighten: 1
1person: 1
namesdalla: 1
borderlin: 1
letsworktogeth: 1
dope: 1
staffcust: 1
lgaord: 1
711: 1
unempathet: 1
fulldont: 1
pri: 1
opsec: 1
leisur: 1
master: 1
companyi: 1
4411: 1
aircargo: 1
bcn: 1
ltltltltlt: 1
whatsoev: 1
ltltltltltlt: 1
correspond: 1
2600: 1
anticonsum: 1
welfar: 1
turquoise: 1
ua1121: 1

```

aa504: 1
aa200: 1
cst: 1
golden: 1
custi: 1
pvr: 1
duck: 1
robertsamps1: 1
stuffi: 1
296: 1
arbitrarili: 1
passengersdont: 1
retribut: 1
3078: 1
cxldprotect: 1
nocharg: 1
nycbueno: 1
john": 1
delays: 1
blackberry10: 1
2450

```

Count Vectorizer

```

In [ ]: cv = CountVectorizer(max_features = 2471)
cv.fit(tweet_list)

```

```

Out[ ]: ▼ CountVectorizer
CountVectorizer(max_features=2471)

```

```

In [ ]: #create bag of words
bow = cv.transform(tweet_list)
bow.shape

#create bag of words dataframe
bow_df = pd.DataFrame(bow.toarray())
bow_df.columns = cv.get_feature_names_out()
bow_df.head()

```

```

Out[ ]: (14640, 2471)

```

```

Out[ ]:
   10  100  1000  10000  1024  1030  105  1051  10pm  11  1130  1130pm  12  1200  1230
0  0    0    0      0      0    0    0    0    0    0    0      0    0    0    0
1  0    0    0      0      0    0    0    0    0    0    0      0    0    0    0
2  0    0    0      0      0    0    0    0    0    0    0      0    0    0    0
3  0    0    0      0      0    0    0    0    0    0    0      0    0    0    0
4  0    0    0      0      0    0    0    0    0    0    0      0    0    0    0

```

Join Dataframes

```
In [ ]: tweets_df2 = pd.concat([tweets_df1, bow_df.set_axis(tweets_df1.index)], axis=1)
tweets_df2.info()
tweets_df2.head()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14640 entries, 0 to 14639
Columns: 2493 entries, tweet_id to zone
dtypes: category(1), datetime64[ns, pytz.FixedOffset(-480)](1), float64(3), in
t64(2475), object(13)
memory usage: 278.4+ MB
```

```
Out[ ]:      tweet_id  airline_sentiment  airline_sentiment_confidence  negativereason  negati
```

	tweet_id	airline_sentiment	airline_sentiment_confidence	negativereason	negati
0	570306133677760513	neutral	1.0000	NaN	

1	570301130888122368	positive	0.3486	NaN	
---	--------------------	----------	--------	-----	--

2	570301083672813571	neutral	0.6837	NaN	
---	--------------------	---------	--------	-----	--

3	570301031407624196	negative	1.0000	Bad Flight	
---	--------------------	----------	--------	------------	--

4	570300817074462722	negative	1.0000	Can't Tell	
---	--------------------	----------	--------	------------	--

Drop Irrelevant Columns

```
In [ ]: tweets_df3 = tweets_df2.copy()
```

```
In [ ]: columns = ['airline_sentiment', 'airline_sentiment_confidence', 'negativereason',
                  'negativereason_confidence', 'name', 'negativereason_gold', 'text',
                  'tweet_created', 'tweet_location', 'user_timezone', 'weighted_sentiment',
                  'unclean_timezone']

tweets_df3 = tweets_df3.drop(columns, axis = 1)
tweets_df3.head()
```

```
Out[ ]:
```

	tweet_id	airline	retweet_count	numeric_sentiment	imputed_neg_reason	clean_timezone
0	570306133677760513	Virgin America	0	0	NaN	(U
1	570301130888122368	Virgin America	0	1	NaN	Paci
2	570301083672813571	Virgin America	0	0	NaN	(U
3	570301031407624196	Virgin America	0	-1	Bad Flight	Paci
4	570300817074462722	Virgin America	0	-1	Can't Tell	Paci

Encoding

```
In [ ]: #encode categorical variables

# Converting type of columns to category
tweets_df3['airline'] = tweets_df3['airline'].astype('category')
tweets_df3['imputed_neg_reason'] = tweets_df3['imputed_neg_reason'].astype('category')
tweets_df3['clean_timezone'] = tweets_df3['clean_timezone'].astype('category')
tweets_df3['time_of_day'] = tweets_df3['time_of_day'].astype('category')

# Assigning numerical values and storing it in another columns
tweets_df3['airline_new'] = tweets_df3['airline'].cat.codes
tweets_df3['imputed_neg_reason_new'] = tweets_df3['imputed_neg_reason'].cat.codes
tweets_df3['clean_timezone_new'] = tweets_df3['clean_timezone'].cat.codes
tweets_df3['time_of_day_new'] = tweets_df3['time_of_day'].cat.codes

# one hot encoder
enc = OneHotEncoder()

# Passing encoded columns
enc_data = pd.DataFrame(enc.fit_transform(
    tweets_df3[['airline_new', 'imputed_neg_reason_new', 'clean_timezone_new',
    tweets_df3['time_of_day_new']].values).toarray().tolist(),
    columns=['airline_new', 'imputed_neg_reason_new', 'clean_timezone_new', 'time_of_day_new'])

# Merge with main
tweets_df4 = tweets_df3.join(enc_data)
tweets_df4.columns = tweets_df4.columns.astype(str)
tweets_df4 = tweets_df4.drop(['airline', 'imputed_neg_reason', 'clean_timezone', 'time_of_day'])

tweets_df4.head()
tweets_df4.info()
```

```
Out[ ]:
```

	tweet_id	retweet_count	numeric_sentiment	tweet_day	10	100	1000	10000
0	570306133677760513	0	0	24	0	0	0	0
1	570301130888122368	0	1	24	0	0	0	0
2	570301083672813571	0	0	24	0	0	0	0
3	570301031407624196	0	-1	24	0	0	0	0
4	570300817074462722	0	-1	24	0	0	0	0

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14640 entries, 0 to 14639
Columns: 2511 entries, tweet_id to 33
dtypes: float64(34), int64(2473), int8(4)
memory usage: 280.1 MB
```

Create Models

```
In [ ]: #split train and test datasets 80/20 split
train_df, rest_df = train_test_split(tweets_df4, train_size=0.8, shuffle=False)

#split validation and test data 10/10 split
val_df, test_df = train_test_split(rest_df, test_size = 0.5, shuffle = False)
```

```
In [ ]: #get targets
train_target = train_df['numeric_sentiment']
train_tweet_ids = train_df['tweet_id']
train_df = train_df.drop(['numeric_sentiment', 'tweet_id'], axis = 1)

val_target = val_df['numeric_sentiment']
val_tweet_ids = val_df['tweet_id']
val_df = val_df.drop(['numeric_sentiment', 'tweet_id'], axis = 1)

test_target = test_df['numeric_sentiment']
test_tweet_ids = test_df['tweet_id']
test_df = test_df.drop(['numeric_sentiment', 'tweet_id'], axis = 1)
```

```
In [ ]: train_df.head()
```

```
Out[ ]:
```

	retweet_count	tweet_day	10	100	1000	10000	1024	1030	105	1051	10pm	11	1130
0	0	24	0	0	0	0	0	0	0	0	0	0	0
1	0	24	0	0	0	0	0	0	0	0	0	0	0
2	0	24	0	0	0	0	0	0	0	0	0	0	0
3	0	24	0	0	0	0	0	0	0	0	0	0	0
4	0	24	0	0	0	0	0	0	0	0	0	0	0

Naive Bayes

```
In [ ]: X = train_df
y = train_target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20, ra
```

```
In [ ]: nb = GaussianNB()
print('Current Parameters:\n')
pprint(nb.get_params())
```

Current Parameters:

```
{'priors': None, 'var_smoothing': 1e-09}
```



```
In [ ]: classifier = GaussianNB()
classifier.fit(X_train, y_train)

cv_scores = cross_val_score(classifier, X, y, cv=5)

print(classifier, ' mean accuracy: ', round(cv_scores.mean()*100, 3), '% std: ')
```

```
Out[ ]: GaussianNB
GaussianNB()

GaussianNB() mean accuracy: 82.608 % std: 0.019 %
```

Tune Hyperparameters

```
In [ ]: np.logspace(0,-9, num=10)

cv_method = RepeatedStratifiedKFold(n_splits=5,
                                     n_repeats=3,
                                     random_state=999)
```

```
Out[ ]: array([1.e+00, 1.e-01, 1.e-02, 1.e-03, 1.e-04, 1.e-05, 1.e-06, 1.e-07,
               1.e-08, 1.e-09])
```

```
In [ ]: params_nb = {'var_smoothing': np.logspace(0,-9, num = 100)}

gcv_nb = GridSearchCV(estimator = classifier,
                      param_grid = params_nb,
                      cv = cv_method,
                      verbose = 1,
                      scoring = 'accuracy')

df_transformed = PowerTransformer().fit_transform(X_test)

gcv_nb.fit(df_transformed, y_test)

model_nb = gcv_nb
```

Fitting 15 folds for each of 100 candidates, totalling 1500 fits

```
Out[ ]: GridSearchCV
        estimator: GaussianNB
        GaussianNB
```

```
In [ ]: gcv_nb.best_score_

gcv_nb.best_params_
```

```
Out[ ]: 0.6975443296338818
```

```
Out[ ]: {'var_smoothing': 0.0012328467394420659}
```

Fit Model

```
In [ ]: y_pred = model_nb.predict(X_test)

cm = confusion_matrix(y_test, y_pred)
print(cm)

accuracy_score(y_test, y_pred)
```

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:486: UserWarning: X has feature names, but GaussianNB was fitted without feature names

```
warnings.warn(
[[1391   6   24]
 [   0 295  206]
 [   0   0 421]])
```

```
Out[ ]: 0.8992744344857021
```

Fit Test Data

```
In [ ]: nb_pred = model_nb.predict(test_df)

cm = confusion_matrix(test_target, nb_pred)
print(cm)

accuracy_score(test_target, nb_pred)
```

```
[[1114   6   7]
 [   0  75 125]
 [   0  25 112]]
```

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:486: UserWarning: X has feature names, but GaussianNB was fitted without feature names

```
warnings.warn(
```

```
Out[ ]: 0.8886612021857924
```

Random Forest Classifier

```
In [ ]: rf = RandomForestClassifier(random_state = 12)
print('Current Parameters:\n')
pprint(nb.get_params())
```

Current Parameters:

```
{'priors': None, 'var_smoothing': 1e-09}
```

```
In [ ]: classifier = RandomForestClassifier(n_estimators = 400, criterion = 'entropy',
classifier.fit(X_train, y_train)
```

```
Out[ ]: ▼ RandomForestClassifier
```

```
RandomForestClassifier(criterion='entropy', n_estimators=400, random_s
tate=12)
```

```
In [ ]: y_pred = classifier.predict(X_test)

cm = confusion_matrix(y_test, y_pred)
print(cm)
```

```
accuracy_score(y_test, y_pred)

print(classification_report(y_pred, y_test))
```

```
[[1421    0    0]
 [   0  439   62]
 [   0  128  293]]
Out[ ]: 0.9189073836961161
```

	precision	recall	f1-score	support
-1	1.00	1.00	1.00	1421
0	0.88	0.77	0.82	567
1	0.70	0.83	0.76	355
accuracy			0.92	2343
macro avg	0.86	0.87	0.86	2343
weighted avg	0.92	0.92	0.92	2343

Tune Hyperparameters

```
In [ ]: param_grid = {
    'n_estimators': [25, 50, 100, 150],
    'max_features': ['sqrt', 'log2', None],
    'max_depth': [3, 6, 9],
    'max_leaf_nodes': [3, 6, 9],
}

grid_search = GridSearchCV(RandomForestClassifier(),
                           param_grid=param_grid)

grid_search.fit(X_train, y_train)
print(grid_search.best_estimator_)
```

```
Out[ ]: GridSearchCV
  estimator: RandomForestClassifier
    RandomForestClassifier
```

RandomForestClassifier(max_depth=9, max_features=None, max_leaf_nodes=9)

Fit Model

```
In [ ]: model_rf = RandomForestClassifier(max_depth = 9,
                                         max_features = None,
                                         max_leaf_nodes = 9,
                                         n_estimators = 25)

model_rf.fit(X_train, y_train)
y_pred_rf = model_rf.predict(X_test)

cm = confusion_matrix(y_test, y_pred)
print(cm)

accuracy_score(y_test, y_pred)
```

```
print(classification_report(y_pred, y_test))
```

Out[]:

```
RandomForestClassifier
RandomForestClassifier(max_depth=9, max_features=None, max_leaf_nodes=
9,
                        n_estimators=25)
```

```
[[1421    0    0]
 [    0  439   62]
 [    0  128  293]]
0.9189073836961161
```

Out[]:

	precision	recall	f1-score	support
-1	1.00	1.00	1.00	1421
0	0.88	0.77	0.82	567
1	0.70	0.83	0.76	355
accuracy			0.92	2343
macro avg	0.86	0.87	0.86	2343
weighted avg	0.92	0.92	0.92	2343

Fit Test Data

In []:

```
rf_pred = model_rf.predict(test_df)

cm = confusion_matrix(test_target, rf_pred)
print(cm)

accuracy_score(test_target, rf_pred)
```

```
[[1127    0    0]
 [    0  171   29]
 [    0   44   93]]
0.950136612021858
```

Out[]:

LSTM Model

In []:

```
tweets_df5 = tweets_df1.copy()
tweets_df5['clean_text'] = pd.Series(tweet_list)
tweets_df5.head()
```

Out []:

	tweet_id	airline_sentiment	airline_sentiment_confidence	negativereason	negati
--	----------	-------------------	------------------------------	----------------	--------

0	570306133677760513	neutral	1.0000	NaN
---	--------------------	---------	--------	-----

1	570301130888122368	positive	0.3486	NaN
---	--------------------	----------	--------	-----

2	570301083672813571	neutral	0.6837	NaN
---	--------------------	---------	--------	-----

3	570301031407624196	negative	1.0000	Bad Flight
---	--------------------	----------	--------	------------

4	570300817074462722	negative	1.0000	Can't Tell
---	--------------------	----------	--------	------------

Tokenize

```
In [ ]: max_fatures = 2471
tokenizer = Tokenizer(num_words=max_fatures, split=' ')
tokenizer.fit_on_texts(tweets_df5['clean_text'].values)
X = tokenizer.texts_to_sequences(tweets_df5['clean_text'].values)
X = pad_sequences(X)
```

Build Model

```
In [ ]: embed_dim = 128
lstm_out = 196

model_lstm = Sequential()
model_lstm.add(Embedding(max_fatures, embed_dim, input_length = X.shape[1]))
model_lstm.add(SpatialDropout1D(0.2))
model_lstm.add(LSTM(lstm_out, dropout=0.2, recurrent_dropout=0.2))
model_lstm.add(Dense(3, activation='softmax'))
model_lstm.compile(loss = 'categorical_crossentropy', optimizer='adam', metrics=
print(model_lstm.summary())
```

```
In [ ]: Y = pd.get_dummies(tweets_df5['numeric_sentiment']).values
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.33, ran
print(X_train.shape, Y_train.shape)
print(X_test.shape, Y_test.shape)
```

Fit Model

```
In [ ]: batch_size = 32
        model_lstm.fit(X_train, Y_train, epochs = 15, batch_size = batch_size, verbose
```

```
Epoch 1/15
307/307 - 42s - loss: 0.6740 - accuracy: 0.7210 - 42s/epoch - 137ms/step
Epoch 2/15
307/307 - 35s - loss: 0.4686 - accuracy: 0.8157 - 35s/epoch - 114ms/step
Epoch 3/15
307/307 - 38s - loss: 0.4006 - accuracy: 0.8427 - 38s/epoch - 124ms/step
Epoch 4/15
307/307 - 35s - loss: 0.3563 - accuracy: 0.8612 - 35s/epoch - 115ms/step
Epoch 5/15
307/307 - 34s - loss: 0.3165 - accuracy: 0.8774 - 34s/epoch - 112ms/step
Epoch 6/15
307/307 - 36s - loss: 0.2749 - accuracy: 0.8894 - 36s/epoch - 117ms/step
Epoch 7/15
307/307 - 34s - loss: 0.2418 - accuracy: 0.9083 - 34s/epoch - 112ms/step
Epoch 8/15
307/307 - 34s - loss: 0.2129 - accuracy: 0.9179 - 34s/epoch - 112ms/step
Epoch 9/15
307/307 - 35s - loss: 0.1930 - accuracy: 0.9275 - 35s/epoch - 114ms/step
Epoch 10/15
307/307 - 33s - loss: 0.1742 - accuracy: 0.9320 - 33s/epoch - 106ms/step
Epoch 11/15
307/307 - 34s - loss: 0.1535 - accuracy: 0.9389 - 34s/epoch - 112ms/step
Epoch 12/15
307/307 - 35s - loss: 0.1415 - accuracy: 0.9458 - 35s/epoch - 115ms/step
Epoch 13/15
307/307 - 36s - loss: 0.1261 - accuracy: 0.9528 - 36s/epoch - 116ms/step
Epoch 14/15
307/307 - 34s - loss: 0.1188 - accuracy: 0.9577 - 34s/epoch - 112ms/step
Epoch 15/15
307/307 - 34s - loss: 0.1112 - accuracy: 0.9587 - 34s/epoch - 112ms/step
Out[ ]: <keras.src.callbacks.History at 0x7fa05cc55cc0>
```

```
In [ ]: validation_size = 1500

X_validate = X_test[:-validation_size]
Y_validate = Y_test[:-validation_size]
X_test = X_test[-validation_size:]
Y_test = Y_test[-validation_size:]
score, acc = model_lstm.evaluate(X_test, Y_test, verbose = 2, batch_size = batch_size)
print("score: %.2f" % (score))
print("acc: %.2f" % (acc))

105/105 - 5s - loss: 1.3055 - accuracy: 0.7473 - 5s/epoch - 47ms/step
score: 1.31
acc: 0.75
```

GRU Model

Tokenize

```
In [ ]: # Preprocessing
        X = tweets_df5["clean_text"]
        y = pd.get_dummies(tweets_df5['numeric_sentiment']).values

        # Tokenization
```

```

tokenizer = Tokenizer()
tokenizer.fit_on_texts(X)
X_seq = tokenizer.texts_to_sequences(X)
vocab_size = len(tokenizer.word_index) + 1

# Padding sequences
max_length = max([len(seq) for seq in X_seq])
X_pad = pad_sequences(X_seq, maxlen=max_length, padding='post')

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X_pad, y, test_size = 0.1,

```

Build Model

```

In [ ]: # Build GRU model
model_gru = Sequential()
model_gru.add(Embedding(vocab_size, 100, input_length=max_length))
model_gru.add(GRU(64))
model_gru.add(Dropout(0.5))
model_gru.add(Dense(3, activation='softmax'))
model_gru.compile(optimizer='adam', loss='categorical_crossentropy', metrics=[
print(model_gru.summary())

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
embedding (Embedding)	(None, 23, 100)	1200400
gru (GRU)	(None, 64)	31872
dropout (Dropout)	(None, 64)	0
dense (Dense)	(None, 3)	195

```

=====
Total params: 1232467 (4.70 MB)
Trainable params: 1232467 (4.70 MB)
Non-trainable params: 0 (0.00 Byte)

```

None

Fit Model

```

In [ ]: batch_size = 64
model_gru.fit(X_train, y_train, epochs=15, batch_size=batch_size, validation_data=

```

```

Epoch 1/15
206/206 [=====] - 11s 40ms/step - loss: 0.8922 - accuracy: 0.6330 - val_loss: 0.6645 - val_accuracy: 0.6967
Epoch 2/15
206/206 [=====] - 7s 36ms/step - loss: 0.5903 - accuracy: 0.7516 - val_loss: 0.5615 - val_accuracy: 0.7760
Epoch 3/15
206/206 [=====] - 7s 35ms/step - loss: 0.4488 - accuracy: 0.8290 - val_loss: 0.5624 - val_accuracy: 0.7937
Epoch 4/15
206/206 [=====] - 10s 50ms/step - loss: 0.3294 - accuracy: 0.8854 - val_loss: 0.5608 - val_accuracy: 0.7951
Epoch 5/15
206/206 [=====] - 10s 48ms/step - loss: 0.2436 - accuracy: 0.9148 - val_loss: 0.5877 - val_accuracy: 0.7828
Epoch 6/15
206/206 [=====] - 7s 33ms/step - loss: 0.1858 - accuracy: 0.9409 - val_loss: 0.7654 - val_accuracy: 0.7746
Epoch 7/15
206/206 [=====] - 8s 38ms/step - loss: 0.1534 - accuracy: 0.9533 - val_loss: 0.8890 - val_accuracy: 0.7616
Epoch 8/15
206/206 [=====] - 9s 42ms/step - loss: 0.1312 - accuracy: 0.9583 - val_loss: 0.8267 - val_accuracy: 0.7657
Epoch 9/15
206/206 [=====] - 7s 36ms/step - loss: 0.1093 - accuracy: 0.9676 - val_loss: 0.9277 - val_accuracy: 0.7650
Epoch 10/15
206/206 [=====] - 8s 39ms/step - loss: 0.1008 - accuracy: 0.9693 - val_loss: 0.9182 - val_accuracy: 0.7555
Epoch 11/15
206/206 [=====] - 7s 34ms/step - loss: 0.0864 - accuracy: 0.9760 - val_loss: 0.9966 - val_accuracy: 0.7555
Epoch 12/15
206/206 [=====] - 8s 38ms/step - loss: 0.0779 - accuracy: 0.9767 - val_loss: 1.0102 - val_accuracy: 0.7589
Epoch 13/15
206/206 [=====] - 7s 36ms/step - loss: 0.0748 - accuracy: 0.9772 - val_loss: 1.0580 - val_accuracy: 0.7616
Epoch 14/15
206/206 [=====] - 7s 35ms/step - loss: 0.0673 - accuracy: 0.9800 - val_loss: 1.0818 - val_accuracy: 0.7548
Epoch 15/15
206/206 [=====] - 8s 37ms/step - loss: 0.0579 - accuracy: 0.9841 - val_loss: 1.2617 - val_accuracy: 0.7575
Out[ ]: <keras.src.callbacks.History at 0x7823c706dff0>

```

```

In [ ]: k_fold = KFold(n_splits=5, shuffle=True, random_state=42)

# Initialize lists to store evaluation metrics
scores = []

# Define callbacks
early_stopping = EarlyStopping(monitor='val_loss', patience=3, restore_best_weights=True)
model_checkpoint = ModelCheckpoint('best_model_cv.h5', monitor='val_loss', save_best_only=True)

# Perform K-fold cross-validation
for fold_index, (train_indices, val_indices) in enumerate(k_fold.split(X_pad, y)):

```



```

print(f"Fold {fold_index + 1}/{k_fold.n_splits}")

# Split data into training and validation sets for this fold
X_train_fold, X_val_fold = X_pad[train_indices], X_pad[val_indices]
y_train_fold, y_val_fold = y[train_indices], y[val_indices]

# Build GRU model
model_gru_cv = Sequential()
model_gru_cv.add(Embedding(vocab_size, 100, input_length=max_length))
model_gru_cv.add(GRU(64))
model_gru_cv.add(Dropout(0.5))
model_gru_cv.add(Dense(3, activation='softmax'))
model_gru_cv.compile(optimizer='adam', loss='categorical_crossentropy', me

# Train the model with callbacks
history = model_gru_cv.fit(X_train_fold, y_train_fold, epochs=15, batch_si:
                           validation_data=(X_val_fold, y_val_fold),
                           callbacks=[early_stopping, model_checkpoint],
                           verbose=1)

# Evaluate the model on the validation set
score = model_gru_cv.evaluate(X_val_fold, y_val_fold, verbose=0)
print("Validation Score:", score)
scores.append(score)

# Calculate mean and standard deviation of validation scores
mean_score = np.mean(scores, axis=0)
std_score = np.std(scores, axis=0)
print("Mean Validation Score:", mean_score)
print("Standard Deviation of Validation Score:", std_score)

```

Fold 1/5
Epoch 1/15
183/183 [=====] - 14s 55ms/step - loss: 0.9304 - accu
racy: 0.6195 - val_loss: 0.7839 - val_accuracy: 0.6452
Epoch 2/15
2/183 [.....] - ETA: 9s - loss: 0.8669 - accuracy:
0.5859

/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103: Use
rWarning: You are saving your model as an HDF5 file via `model.save()`. This f
ile format is considered legacy. We recommend using instead the native Keras f
ormat, e.g. `model.save('my_model.keras')`.
saving_api.save_model(

183/183 [=====] - 7s 38ms/step - loss: 0.6474 - accuracy: 0.7228 - val_loss: 0.5552 - val_accuracy: 0.7790
Epoch 3/15
183/183 [=====] - 7s 39ms/step - loss: 0.4621 - accuracy: 0.8233 - val_loss: 0.5903 - val_accuracy: 0.7913
Epoch 4/15
183/183 [=====] - 6s 33ms/step - loss: 0.3428 - accuracy: 0.8774 - val_loss: 0.6059 - val_accuracy: 0.7753
Epoch 5/15
183/183 [=====] - 6s 33ms/step - loss: 0.2488 - accuracy: 0.9166 - val_loss: 0.6056 - val_accuracy: 0.7879
Validation Score: [0.5552219152450562, 0.7790300250053406]
Fold 2/5
Epoch 1/15
183/183 [=====] - 9s 38ms/step - loss: 0.9262 - accuracy: 0.6254 - val_loss: 0.8963 - val_accuracy: 0.6253
Epoch 2/15
183/183 [=====] - 6s 31ms/step - loss: 0.6521 - accuracy: 0.7180 - val_loss: 0.5901 - val_accuracy: 0.7599
Epoch 3/15
183/183 [=====] - 7s 36ms/step - loss: 0.4700 - accuracy: 0.8114 - val_loss: 0.5878 - val_accuracy: 0.7678
Epoch 4/15
183/183 [=====] - 6s 32ms/step - loss: 0.3530 - accuracy: 0.8742 - val_loss: 0.6337 - val_accuracy: 0.7637
Epoch 5/15
183/183 [=====] - 7s 36ms/step - loss: 0.2540 - accuracy: 0.9150 - val_loss: 0.6880 - val_accuracy: 0.7572
Epoch 6/15
183/183 [=====] - 6s 32ms/step - loss: 0.1854 - accuracy: 0.9415 - val_loss: 0.8130 - val_accuracy: 0.7374
Validation Score: [0.5878012180328369, 0.7677595615386963]
Fold 3/5
Epoch 1/15
183/183 [=====] - 10s 41ms/step - loss: 0.9252 - accuracy: 0.6249 - val_loss: 0.7838 - val_accuracy: 0.6332
Epoch 2/15
183/183 [=====] - 6s 34ms/step - loss: 0.6305 - accuracy: 0.7285 - val_loss: 0.5670 - val_accuracy: 0.7766
Epoch 3/15
183/183 [=====] - 8s 44ms/step - loss: 0.4694 - accuracy: 0.8204 - val_loss: 0.5755 - val_accuracy: 0.7715
Epoch 4/15
183/183 [=====] - 7s 37ms/step - loss: 0.3610 - accuracy: 0.8690 - val_loss: 0.5919 - val_accuracy: 0.7859
Epoch 5/15
183/183 [=====] - 6s 33ms/step - loss: 0.2594 - accuracy: 0.9147 - val_loss: 0.6682 - val_accuracy: 0.7773
Validation Score: [0.5670248866081238, 0.7766393423080444]
Fold 4/5
Epoch 1/15
183/183 [=====] - 10s 42ms/step - loss: 0.9033 - accuracy: 0.6323 - val_loss: 0.7586 - val_accuracy: 0.6643
Epoch 2/15
183/183 [=====] - 6s 34ms/step - loss: 0.6010 - accuracy: 0.7459 - val_loss: 0.6046 - val_accuracy: 0.7490
Epoch 3/15
183/183 [=====] - 7s 36ms/step - loss: 0.4442 - accuracy: 0.8323 - val_loss: 0.5934 - val_accuracy: 0.7572
Epoch 4/15

```

183/183 [=====] - 8s 41ms/step - loss: 0.3351 - accuracy: 0.8850 - val_loss: 0.6892 - val_accuracy: 0.7602
Epoch 5/15
183/183 [=====] - 6s 35ms/step - loss: 0.2530 - accuracy: 0.9159 - val_loss: 0.7233 - val_accuracy: 0.7520
Epoch 6/15
183/183 [=====] - 7s 38ms/step - loss: 0.1976 - accuracy: 0.9378 - val_loss: 0.8210 - val_accuracy: 0.7459
Validation Score: [0.5934033393859863, 0.7571721076965332]
Fold 5/5
Epoch 1/15
183/183 [=====] - 10s 40ms/step - loss: 0.9287 - accuracy: 0.6242 - val_loss: 0.8679 - val_accuracy: 0.6305
Epoch 2/15
183/183 [=====] - 8s 42ms/step - loss: 0.6468 - accuracy: 0.7183 - val_loss: 0.5907 - val_accuracy: 0.7579
Epoch 3/15
183/183 [=====] - 6s 35ms/step - loss: 0.4668 - accuracy: 0.8173 - val_loss: 0.5861 - val_accuracy: 0.7681
Epoch 4/15
183/183 [=====] - 6s 35ms/step - loss: 0.3529 - accuracy: 0.8691 - val_loss: 0.6392 - val_accuracy: 0.7824
Epoch 5/15
183/183 [=====] - 6s 34ms/step - loss: 0.2570 - accuracy: 0.9126 - val_loss: 0.6333 - val_accuracy: 0.7684
Epoch 6/15
183/183 [=====] - 7s 36ms/step - loss: 0.1924 - accuracy: 0.9388 - val_loss: 0.7228 - val_accuracy: 0.7674
Validation Score: [0.5861008167266846, 0.7681010961532593]
Mean Validation Score: [0.57791044 0.76974043]
Standard Deviation of Validation Score: [0.01440975 0.00772613]

```

```
In [ ]: history_final = model_gru.fit(X_pad, y, epochs=15, batch_size=batch_size, verbose=0)
```

```

Epoch 1/15
229/229 [=====] - 17s 75ms/step - loss: 0.1504 - accu
racy: 0.9616
Epoch 2/15
229/229 [=====] - 13s 56ms/step - loss: 0.0961 - accu
racy: 0.9733
Epoch 3/15
229/229 [=====] - 13s 58ms/step - loss: 0.0720 - accu
racy: 0.9803
Epoch 4/15
229/229 [=====] - 13s 59ms/step - loss: 0.0616 - accu
racy: 0.9823
Epoch 5/15
229/229 [=====] - 11s 49ms/step - loss: 0.0587 - accu
racy: 0.9838
Epoch 6/15
229/229 [=====] - 8s 36ms/step - loss: 0.0496 - accur
acy: 0.9859
Epoch 7/15
229/229 [=====] - 13s 58ms/step - loss: 0.0498 - accu
racy: 0.9857
Epoch 8/15
229/229 [=====] - 13s 57ms/step - loss: 0.0476 - accu
racy: 0.9861
Epoch 9/15
229/229 [=====] - 13s 57ms/step - loss: 0.0474 - accu
racy: 0.9859
Epoch 10/15
229/229 [=====] - 8s 33ms/step - loss: 0.0485 - accur
acy: 0.9850
Epoch 11/15
229/229 [=====] - 9s 37ms/step - loss: 0.0445 - accur
acy: 0.9859
Epoch 12/15
229/229 [=====] - 8s 37ms/step - loss: 0.0409 - accur
acy: 0.9874
Epoch 13/15
229/229 [=====] - 7s 32ms/step - loss: 0.0394 - accur
acy: 0.9876
Epoch 14/15
229/229 [=====] - 8s 37ms/step - loss: 0.0361 - accur
acy: 0.9885
Epoch 15/15
229/229 [=====] - 8s 35ms/step - loss: 0.0331 - accur
acy: 0.9896

```

```

In [ ]: # Predict probabilities for each class on the entire dataset
        probs = model_gru.predict(X_pad, verbose=2, batch_size=batch_size)

        # Determine predicted classes based on the highest probability
        y_pred = np.argmax(probs, axis=1)

        # Convert one-hot encoded labels to single label for the entire dataset
        y_true = np.argmax(y, axis=1)

        # Generate classification report for the entire dataset
        print("Classification Report:")
        print(classification_report(y_true, y_pred))

```

229/229 - 2s - 2s/epoch - 10ms/step

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	9178
1	0.98	0.98	0.98	3099
2	0.98	0.99	0.98	2363
accuracy			0.99	14640
macro avg	0.99	0.99	0.99	14640
weighted avg	0.99	0.99	0.99	14640